


# **Fine-tuning BERT for Text Classification**

**Overview with Example Code**

**Shaw Talebi**



# Fine-tuning BERT for Text Classification

1

2

3

# Fine-tuning

Adapting a pre-trained model to a particular task (through additional training)



**Pre-trained Model**

(Self-supervised)

***~3B Words***

(BERT)

**Fine-tuned Model**

(Supervised)

***~3k Examples***

(MRCP)

# Fine-tuning

Adapting a pre-trained model to a particular task (through additional training)



**Fine-tuned  
Model**


**Additional  
Fine-tuning**



# BERT

Language model developed for fine-tuning

## Task #1: Masked LM




# BERT

Language model developed for fine-tuning

## Task #2: Next Sentence Prediction

**[CLS]** BERT is conceptually simple and empirically powerful. **[SEP]** It obtains new state-of-the-art results on 11 NLP tasks **[SEP]**





# BERT

Language model developed for fine-tuning

## Task #2: Next Sentence Prediction

**[CLS]** BERT is conceptually simple and empirically powerful. **[SEP]** The cat sat on the mat **[SEP]**



*NotNext*




# Text Classification

Assigning a label to text sequences



Spam Detection




Tier 3

Tier 2

Tier 1

Categorizing IT Tickets



Review Sentiment  
Analysis




# Example Code: Fine-tuning BERT for Phishing URL Identification



Epoch	Training Loss	Validation Loss	Accuracy
1	No log	0.423429	{'accuracy': 0.876}
2	0.412400	0.551401	{'accuracy': 0.878}
3	0.412400	0.593060	{'accuracy': 0.899}
4	0.211600	0.640572	{'accuracy': 0.894}
5	0.211600	0.839775	{'accuracy': 0.891}
6	0.064300	0.930830	{'accuracy': 0.887}
7	0.064300	0.988515	{'accuracy': 0.889}
8	0.020000	1.007572	{'accuracy': 0.884}
9	0.020000	1.040836	{'accuracy': 0.888}
10	0.009500	1.034907	{'accuracy': 0.896}

Overfitting




# Example Code: Fine-tuning BERT for Phishing URL Identification

## Imports

```
from datasets import DatasetDict, Dataset
from transformers import AutoTokenizer, AutoModelForSequenceClassification,
                        TrainingArguments, Trainer

import evaluate
import numpy as np
from transformers import DataCollatorWithPadding
```



## Load Data

```
dataset_dict = load_dataset("shawhin/phishing-site-classification")
```

[\[https://huggingface.co/datasets/shawhin/phishing-site-classification\]](https://huggingface.co/datasets/shawhin/phishing-site-classification)



# Example Code: Fine-tuning BERT for Phishing URL Identification

## Load Pre-trained Model

```
# define pre-trained model path
model_path = "google-bert/bert-base-uncased"

# load model tokenizer
tokenizer = AutoTokenizer.from_pretrained(model_path)

# load model with binary classification head
id2label = {0: "Safe", 1: "Not Safe"}
label2id = {"Safe": 0, "Not Safe": 1}
model = AutoModelForSequenceClassification.from_pretrained(model_path,
                                                         num_labels=2,
                                                         id2label=id2label,
                                                         label2id=label2id,)
```

# Example Code: Fine-tuning BERT for Phishing URL Identification

## Set Trainable Parameters

```
# freeze all base model parameters
for name, param in model.base_model.named_parameters():
    param.requires_grad = False

# unfreeze base model pooling layers
for name, param in model.base_model.named_parameters():
    if "pooler" in name:
        param.requires_grad = True
```



# Example Code: Fine-tuning BERT for Phishing URL Identification

## Data Pre-processing

```
# define text preprocessing
def preprocess_function(examples):
    # return tokenized text with truncation
    return tokenizer(examples["text"], truncation=True)

# preprocess all datasets
tokenized_data = dataset_dict.map(preprocess_function, batched=True)
```

```
# create data collator
data_collator = DataCollatorWithPadding(tokenizer=tokenizer)
```

# Example Code: Fine-tuning BERT for Phishing URL Identification

## Define Evaluation Metrics

```
# load metrics
accuracy = evaluate.load("accuracy")
auc_score = evaluate.load("roc_auc")

def compute_metrics(eval_pred):
    # get predictions
    predictions, labels = eval_pred

    # apply softmax to get probabilities
    probabilities = np.exp(predictions) / np.exp(predictions).sum(-1,
                                                                    keepdims=True)

    # use probabilities of the positive class for ROC AUC
    positive_class_probs = probabilities[:, 1]

    # compute auc
    auc = np.round(auc_score.compute(prediction_scores=positive_class_probs,
                                     references=labels)['roc_auc'], 3)

    # predict most probable class
    predicted_classes = np.argmax(predictions, axis=1)

    # compute accuracy
    acc = np.round(accuracy.compute(predictions=predicted_classes,
                                     references=labels)['accuracy'], 3)

    return {"Accuracy": acc, "AUC": auc}
```



# Example Code: Fine-tuning BERT for Phishing URL Identification

## Training Parameters

```
# hyperparameters
lr = 2e-4
batch_size = 8
num_epochs = 10

training_args = TrainingArguments(
    output_dir="bert-phishing-classifier_teacher",
    learning_rate=lr,
    per_device_train_batch_size=batch_size,
    per_device_eval_batch_size=batch_size,
    num_train_epochs=num_epochs,
    logging_strategy="epoch",
    eval_strategy="epoch",
    save_strategy="epoch",
    load_best_model_at_end=True,
)
```

# Example Code: Fine-tuning BERT for Phishing URL Identification

## Fine-tune Model

```
trainer = Trainer(  
    model=model,  
    args=training_args,  
    train_dataset=tokenized_data["train"],  
    eval_dataset=tokenized_data["test"],  
    tokenizer=tokenizer,  
    data_collator=data_collator,  
    compute_metrics=compute_metrics,  
)  
  
trainer.train()
```

Training Loss	Epoch	Step	Validation Loss	Accuracy	Auc
0.4916	1.0	263	0.4228	0.784	0.915
0.3894	2.0	526	0.3586	0.818	0.932
0.3837	3.0	789	0.3144	0.86	0.939
0.3574	4.0	1052	0.4494	0.807	0.942
0.3517	5.0	1315	0.3287	0.86	0.947
0.3518	6.0	1578	0.3042	0.871	0.949
0.3185	7.0	1841	0.2900	0.862	0.949
0.3267	8.0	2104	0.2958	0.876	0.95
0.3153	9.0	2367	0.2881	0.867	0.951
0.3061	10.0	2630	0.2963	0.873	0.951



# Example Code: Fine-tuning BERT for Phishing URL Identification

## Validation Data

```
# apply model to validation dataset
predictions = trainer.predict(tokenized_data["validation"])


# Extract the logits and labels from the predictions object
logits = predictions.predictions
labels = predictions.label_ids

# Use your compute_metrics function
metrics = compute_metrics((logits, labels))
print(metrics)

# >> {'Accuracy': 0.889, 'AUC': 0.946}
```

**Make LLMs**

**Smaller**



# Fine-tuning BERT for Text Classification

A Hackable Example with Python Code

Although today's 100B+ parameter transformer models are state-of-the-art in AI, there's still much we can accomplish with smaller (< 1B parameter) models. In this article, I will walk through one such example, fine-tuning BERT (110M parameters) to classify phishing URLs. I'll start by covering key concepts and then share example Python code.

**Friend Link in Description**

